



Information gain from conditioning — KL divergence vs naive baseline


Δx state = 0



Δx state = 1



Δx state = 2



KL(cond || naive)

1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00